

Meta-Learning | The MAML method

(1st Coding Lab – DL2025)

Why Meta-Learning?

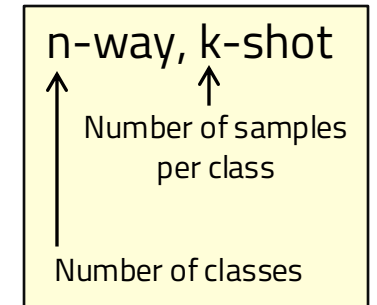
- Humans can generalize from very few samples



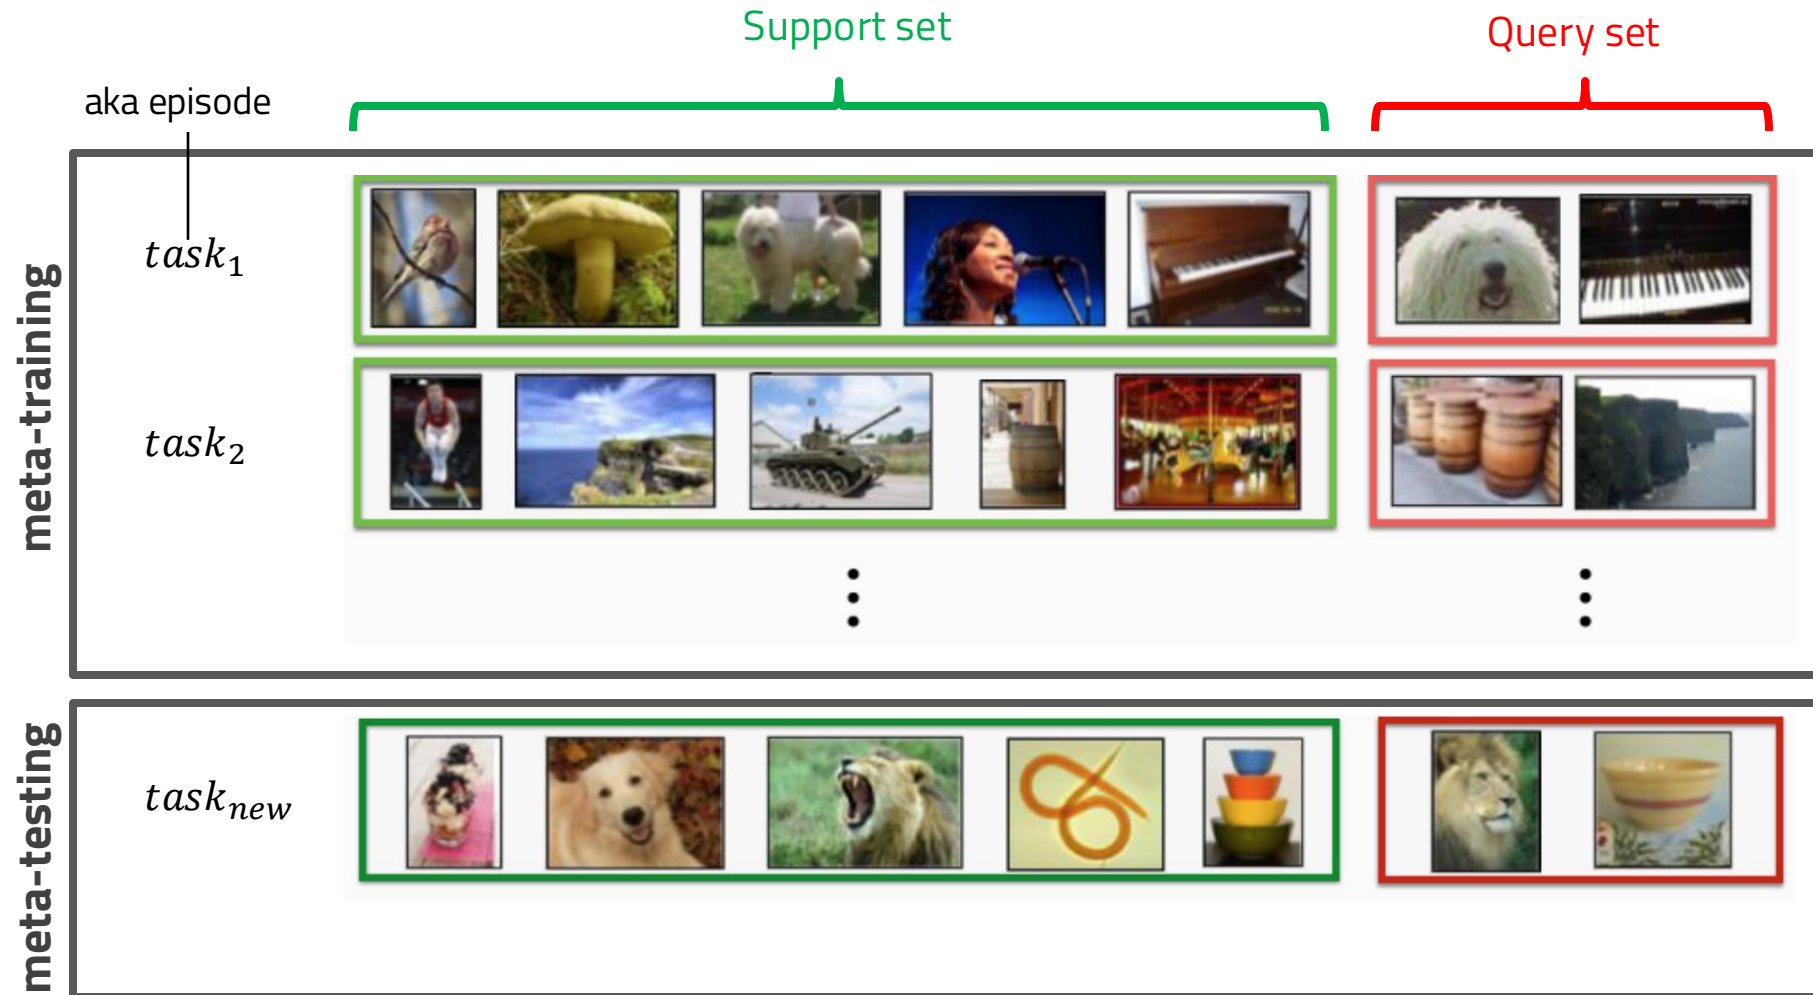
Only one image is enough
for a kid to be able to
classify deers in the future!

- Is this the case for deep neural networks?
- Meta-learning: Methods that generalize from very few samples (e.g., 1 sample per class)
 - Model based (e.g., Hyper Networks)
 - Optimization based (e.g., MAML)
 - Metric based - aka Metric learning (e.g., Matching Nets)

Image classification meta-learning tasks





Meta- training / testing phases



- MAML: **M**odel **A**gnostic **M**eta **L**earning
- An optimization based meta-learning method.

1) Copy model per task 2) Support set train 3) Calculate query set loss 4) Sum task losses



MAML-algorithm

Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
- 2: **while** not done **do**
- 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
- 4: **for all** \mathcal{T}_i **do**
- 5: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples
- 6: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
- 7: **end for**
- 8: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$
- 9: **end while**

Inner task specific model updates


Outer meta-model updates

This gradient depends on the task specific model parameters:

$$\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$$

Which in turn depend on the per task gradient computation.

Computation of the Hessian vector products (2nd derivative)



PyTorch Background:

1. Tensor Operations & Back Propagation
2. Stateful vs. Stateless Models


Meta-Learning Specific:

1. Task Family: Sine waves task distribution.
2. Meta-training with MAML.
3. Model Adaptation: Evaluating MAML against supervised pre-training.

Tensor Operations & Back Propagation

- What is a computational graph?

$$\mathcal{F}(x_1, x_2) = (x_1 + x_2)(x_2 + 1)$$



Forward Pass:

$$v_1 = x_1 + x_2$$

$$v_2 = x_2 + 1$$

$$v_3 = v_1 + v_2$$

Backward Pass:

$$\frac{\partial \mathcal{F}}{\partial v_3} = \frac{\partial v_3}{\partial v_3} = 1$$

$$\frac{\partial \mathcal{F}}{\partial v_1} = \frac{\partial \mathcal{F}}{\partial v_3} \frac{\partial v_3}{\partial v_1} = v_2$$


$$\frac{\partial \mathcal{F}}{\partial v_2} = \frac{\partial \mathcal{F}}{\partial v_3} \frac{\partial v_3}{\partial v_2} = v_1$$

$$\frac{\partial \mathcal{F}}{\partial x_1} = \frac{\partial \mathcal{F}}{\partial v_1} \frac{\partial v_1}{\partial x_1} = v_2$$

$$\frac{\partial \mathcal{F}}{\partial x_2} = \frac{\partial \mathcal{F}}{\partial v_1} \frac{\partial v_1}{\partial x_2} + \frac{\partial \mathcal{F}}{\partial v_2} \frac{\partial v_2}{\partial x_2} = v_1 + v_2$$

Meta-learning regression setting

- Regression task
 - The model has to learn the mapping:
 $F: X \rightarrow Y$
 - The Mean Squared Error (MSE) loss is used for supervision.
- Tasks distribution: $y = \text{amplitude} * \sin(\text{phase} + x)$
 - A family of sines defined by: (amplitude_min, amplitude_max, phase_min, phase_max, x_min, x_max)
- Single task
 - Support set: The x and y (target) coordinates (in blue).
 - Query set: Only the x coordinates (in red).



- MAML paper [[here](#)]
- Blog on pytorch's *create_graph* option for 2nd order derivatives [[here](#)]
- Video tutorials:
 - Meta-learning (the general setting)
 - CS 182: Lecture 21: Part 1: Meta-Learning (UC Berkeley) [[here](#)]
 - MAML (intuitive description) [[here](#)]